

# Throwing Tools at Ranges

Tina Ulbrich







# Throwing Tools at Ranges

Tina Ulbrich – ROSEN Technology and Research Center  
C++ on Sea 2023





# Outline

---

- Setup
- Code example
- Google Benchmark
- Cachegrind
- Visual Studio Profiler
- Optimizing
- Conclusion



Setup

# Setup

## Laptop

Processor	Intel Core i3-6157U CPU @ 2,40GHz
Cores	2
Logical Processors	4
L1-Cache	128 KB
L2-Cache	512 KB
L3-Cache	3 MB
RAM	6 GB
OS	Windows 10

## PC

Processor	13th Gen Intel Core i7-13700KF @ 3,4GHz
Cores	16
Logical Processors	24
L1-Cache	1,4 MB
L2-Cache	24 MB
L3-Cache	30 MB
RAM	64 GB
OS	Windows 11





Code



# Function 1, C++23

```
auto random_23_1(const std::vector<double>& rng)
{
    return ranges::views::cartesian_product(
        rng,
        ranges::views::iota(1, 5),
        rng | ranges::views::transform([](const auto e) { return e * e + 2; })
        | ranges::views::transform([](const auto t)
        {
            const auto& [a, b, c] = t;
            return a * b + 2 * c;
        }));
}
```

# Function 1, C++17

```
auto random_17_1(const std::vector<double>& rng)
{
    auto out = std::vector<double>();
    out.reserve(rng.size() * rng.size() * 4);
    for (const auto e1 : rng)
    {
        for (const auto e2 : std::array{ 1, 2, 3, 4 })
        {
            for (const auto e3 : rng)
            {
                const auto e = e1 * e2 + 2 * (e3 * e3 + 2);
                out.push_back(e);
            }
        }
    }

    return out;
}
```



# Function 1, C – style C++

```
double* random_c_style_1(double* rng, int n)
{
    double* out = new double[n * n * 4];
    int out_idx = 0;
    for (int i = 0; i < n; ++i)
    {
        for (int j = 1; j < 5; ++j)
        {
            for (int k = 0; k < n; ++k)
            {
                out[out_idx] = rng[i] * j + 2 * (rng[k] * rng[k] + 2);
                ++out_idx;
            }
        }
    }

    return out;
}
```

# Function 2, C++23

```
auto random_23_2(auto& rng1)
{
    return ranges::views::zip_with(
        [](const auto a, const auto b) { return a / static_cast<double>(b); },
        rng1,
        ranges::views::iota(10)
    );
}
```



# Function 2, C++17

```
auto random_17_2(const std::vector<double>& rng)
{
    auto out = std::vector<double>(rng.size());
    int i = 10;
    std::transform(rng.begin(), rng.end(), out.begin(), [&](const auto e)
    {
        return e / static_cast<double>(i++);
    });

    return out;
}
```

# Function 2, C – style C++

```
double* random_c_style_2(double* rng, int n)
{
    double* out = new double[n];
    for (int i = 0; i < n; ++i)
    {
        out[i] = rng[i] / double(i + 10);
    }

    return out;
}
```



# Function 3, C++23

```
auto random_23_3(auto& rng1, auto& rng2)
{
    auto tmp1 = rng1 | ranges::views::partial_sum;
    auto tmp2 = rng2
        | ranges::views::sliding(2)
        | ranges::views::transform([](const auto subrng)
        {
            return 3.0 * subrng[1] - 2.0 * subrng[0];
        });

    return ranges::inner_product(tmp1 | ranges::views::drop(1), tmp2, 0.0);
}
```

# Function 3, C++17

```
auto random_17_3(const std::vector<double>& rng1, const std::vector<double>& rng2)
{
    auto tmp1 = std::vector<double>(rng1.size());
    std::partial_sum(rng1.begin(), rng1.end(), tmp1.begin());

    auto tmp2 = std::vector<double>(rng2.size());
    std::adjacent_difference(rng2.begin(), rng2.end(), tmp2.begin(),
                             [] (const auto left, const auto right)
    {
        return 3.0 * left - 2.0 * right;
    });

    return std::inner_product(tmp1.begin() + 1, tmp1.end(), tmp2.begin() + 1, 0.0);
}
```



# Function 3, C – style C++

```
double random_c_style_3(double* rng1, double* rng2, int n)
{
    double* tmp1 = new double[n];
    tmp1[0] = rng1[0];
    for (int i = 1; i < n; ++i)
    {
        tmp1[i] = tmp1[i - 1] + rng1[i];
    }

    double* tmp2 = new double[n - 1];
    for (int i = 0; i < n - 1; ++i)
    {
        tmp2[i] = 3.0 * rng2[i + 1] - 2.0 * rng2[i];
    }

    ...
}
```

# Function 3, C – style C++

```
...  
    double sum = 0.0;  
    for (int i = 0; i < n - 1; ++i)  
    {  
        sum += tmp1[i + 1] * tmp2[i];  
    }  
  
    delete[] tmp1;  
    delete[] tmp2;  
    return sum;  
}
```



# Function 4, C++23

```
double random_23(const std::vector<double>& rng)
{
    auto rng1 = random_23_1(rng);
    auto rng2 = random_23_2(rng1);

    return random_23_3(rng1, rng2);
}
```

# Function 4, C++17

```
double random_17(const std::vector<double>& rng)
{
    const auto rng1 = random_17_1(rng);
    const auto rng2 = random_17_2(rng1);

    return random_17_3(rng1, rng2);
}
```



# Function 4, C – style C++

```
double random_c_style(double* rng, int n)
{
    auto tmp1 = random_c_style_1(rng, n);
    auto tmp2 = random_c_style_2(tmp1, n * n * 4);
    auto out = random_c_style_3(tmp1, tmp2, n * n * 4);

    delete[] tmp1;
    delete[] tmp2;

    return out;
}
```



# Google Benchmark



# Google Benchmark

```
static void cpp_17_benchmark_optimized(benchmark::State& state)
{
    const auto rng = std::vector(data.begin(), data.end());
    for (auto _ : state)
    {
        benchmark::DoNotOptimize(random_17_optimized(rng));
    }
}

BENCHMARK(cpp_17_benchmark)->Unit(benchmark::kMillisecond);

BENCHMARK_MAIN();
```

# Google Benchmark

Run on (24 X 3450.41 MHz CPU s)

CPU Caches:

L1 Data 48 KiB (x12)

L1 Instruction 32 KiB (x12)

L2 Unified 2048 KiB (x12)

L3 Unified 30720 KiB (x1)

Benchmark	Time	CPU	Iterations
c_style_benchmark	25.3 ms	19.8 ms	34
cpp_17_benchmark	35.4 ms	23.8 ms	50
cpp_23_benchmark	317.0 ms	223.0 ms	4





Cachegrind

# Cachegrind, C – style C++

I refs: 92,364,114  
I1 misses: 2,025  
LLi misses: 1,956  
I1 miss rate: 0.00%  
LLi miss rate: 0.00%

D refs: 26,758,586 ( 16,556,648 rd + 10,201,938 wr)  
D1 misses: 4,515,098 ( 2,512,546 rd + 2,002,552 wr)  
LLd misses: 4,509,814 ( 2,508,034 rd + 2,001,780 wr)  
D1 miss rate: 16.9% ( 15.2% + 19.6% )  
LLd miss rate: 16.9% ( 15.1% + 19.6% )

LL refs: 4,517,123 ( 2,514,571 rd + 2,002,552 wr)  
LL misses: 4,511,770 ( 2,509,990 rd + 2,001,780 wr)  
LL miss rate: 3.8% ( 2.3% + 19.6% )

Branches: 11,334,889 ( 11,329,459 cond + 5,430 ind)  
Mispredicts: 18,731 ( 18,163 cond + 568 ind)  
Mispred rate: 0.2% ( 0.2% + 10.5% )



# Cachegrind, C++17

I refs: 252,376,633  
I1 misses: 2,044  
LLi misses: 1,976  
I1 miss rate: 0.00%  
LLi miss rate: 0.00%

D refs: 144,802,262 ( 26,594,447 rd + 118,207,815 wr)  
D1 misses: 6,015,098 ( 2,512,549 rd + 3,502,549 wr)  
LLd misses: 5,979,557 ( 2,508,074 rd + 3,471,483 wr)  
D1 miss rate: 4.2% ( 9.4% + 3.0% )  
LLd miss rate: 4.1% ( 9.4% + 2.9% )

LL refs: 6,017,142 ( 2,514,593 rd + 3,502,549 wr)  
LL misses: 5,981,533 ( 2,510,050 rd + 3,471,483 wr)  
LL miss rate: 1.5% ( 0.9% + 2.9% )

Branches: 115,330,825 (115,325,399 cond + 5,426 ind)  
Mispredicts: 18,725 ( 18,152 cond + 573 ind)  
Mispred rate: 0.0% ( 0.0% + 10.6% )



# Cachegrind, C++23

I refs: 542,456,846  
I1 misses: 2,036  
LLi misses: 1,936  
I1 miss rate: 0.00%  
LLi miss rate: 0.00%

D refs: 232,780,702 (112,575,063 rd + 120,205,639 wr)  
D1 misses: 14,845 ( 12,349 rd + 2,496 wr)  
LLd misses: 9,425 ( 7,708 rd + 1,717 wr)  
D1 miss rate: 0.0% ( 0.0% + 0.0% )  
LLd miss rate: 0.0% ( 0.0% + 0.0% )

LL refs: 16,881 ( 14,385 rd + 2,496 wr)  
LL misses: 11,361 ( 9,644 rd + 1,717 wr)  
LL miss rate: 0.0% ( 0.0% + 0.0% )

Branches: 36,340,555 ( 36,335,152 cond + 5,403 ind)  
Mispredicts: 19,671 ( 19,108 cond + 563 ind)  
Mispred rate: 0.1% ( 0.1% + 10.4% )

# Cachegrind compare

	C-Style	C++17	C++23
I refs:	92,364,114	252,376,633	542,456,846
I1 misses:	2,025	2,044	2,036
LLi misses:	1,956	1,976	1,936
I1 miss rate:	0.00%	0.00%	0.00%
LLi miss rate:	0.00%	0.00%	0.00%
D refs:	26,758,586	144,802,262	232,780,702
D1 misses:	4,515,098	6,015,098	14,845
LLd misses:	4,509,814	5,979,557	9,425
D1 miss rate:	16.9%	4.2%	0.0%
LLd miss rate:	16.9%	4.1%	0.0%
LL refs:	4,517,123	6,017,142	16,881
LL misses:	4,511,770	5,981,533	11,361
LL miss rate:	3.8%	1.5%	0.0%
Branches:	11,334,889	115,330,825	36,340,555
Mispredicts:	18,731	18,725	19,671
Mispred rate:	0.2%	0.0%	0.1%

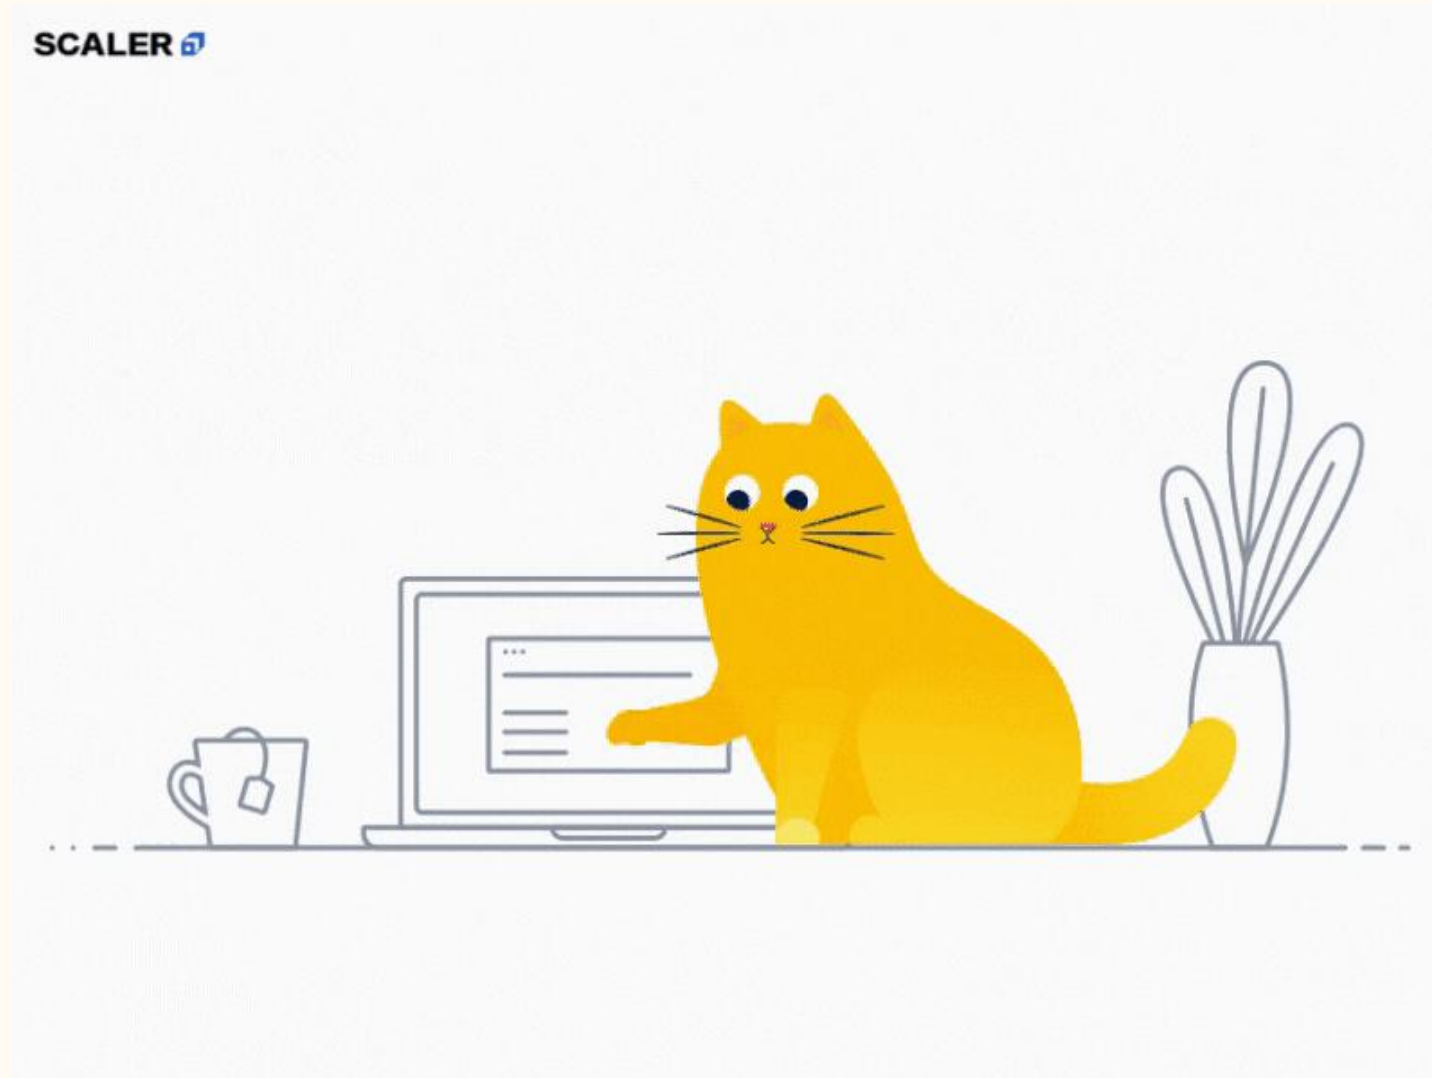




# Profiling and Optimizing



# Visual Studio Profiler



# Function 3, C – style C++

```
double random_c_style_3(double* rng1, double* rng2, int n) {  
    double* tmp1 = new double[n];  
    tmp1[0] = rng1[0];  
    for (int i = 1; i < n; ++i) {  
        tmp1[i] = tmp1[i - 1] + rng1[i];  
    }  
  
    double* tmp2 = new double[n - 1];  
    for (int i = 0; i < n - 1; ++i) {  
        tmp2[i] = 3.0 * rng2[i + 1] - 2.0 * rng2[i];  
    }  
  
    double sum = 0.0;  
    for (int i = 0; i < n - 1; ++i) {  
        sum += tmp1[i + 1] * tmp2[i];  
    }  
  
    delete[] tmp1;  
    delete[] tmp2;  
    return sum;  
}
```

# Optimizing, C – style C++

```
double random_c_style_3(double* rng1, double* rng2, int n)
{
    double partial_sum = rng1[0];
    double sum = 0.0;
    for (int i = 1; i < n; ++i)
    {
        auto partial_sum += rng1[i];
        sum += partial_sum * (3.0 * rng2[i] - 2.0 * rng2[i - 1]);
    }

    return sum;
}
```



# Function 3, C – style C++

```
double random_c_style_3(double* rng1, double* rng2, int n) {  
    double* tmp1 = new double[n];  
    tmp1[0] = rng1[0];  
    for (int i = 1; i < n; ++i) {  
        tmp1[i] = tmp1[i - 1] + rng1[i];  
    }  
  
    double* tmp2 = new double[n - 1];  
    for (int i = 0; i < n - 1; ++i) {  
        tmp2[i] = 3.0 * rng2[i + 1] - 2.0 * rng2[i];  
    }  
  
    double sum = 0.0;  
    for (int i = 0; i < n - 1; ++i) {  
        sum += tmp1[i + 1] * tmp2[i];  
    }  
  
    delete[] tmp1;  
    delete[] tmp2;  
    return sum;  
}
```

# Optimizing, C – style C++

```
double random_c_style_3(double* rng1, double* rng2, int n)
{
    double partial_sum = rng1[0];
    double sum = 0.0;
    for (int i = 1; i < n; ++i)
    {
        auto partial_sum += rng1[i];
        sum += partial_sum * (3.0 * rng2[i] - 2.0 * rng2[i - 1]);
    }

    return sum;
}
```

# Function 3, C – style C++

```
double random_c_style_3(double* rng1, double* rng2, int n) {  
    double* tmp1 = new double[n];  
    tmp1[0] = rng1[0];  
    for (int i = 1; i < n; ++i) {  
        tmp1[i] = tmp1[i - 1] + rng1[i];  
    }  
  
    double* tmp2 = new double[n - 1];  
    for (int i = 0; i < n - 1; ++i) {  
        tmp2[i] = 3.0 * rng2[i + 1] - 2.0 * rng2[i];  
    }  
  
    double sum = 0.0;  
    for (int i = 0; i < n - 1; ++i) {  
        sum += tmp1[i + 1] * tmp2[i];  
    }  
  
    delete[] tmp1;  
    delete[] tmp2;  
    return sum;  
}
```



# Optimizing, C – style C++

```
double random_c_style_3(double* rng1, double* rng2, int n)
{
    double partial_sum = rng1[0];
    double sum = 0.0;
    for (int i = 1; i < n; ++i)
    {
        auto partial_sum += rng1[i];
        sum += partial_sum * (3.0 * rng2[i] - 2.0 * rng2[i - 1]);
    }

    return sum;
}
```

# Function 3, C++17

```
auto random_17_3(const std::vector<double>& rng1, const std::vector<double>& rng2)
{
    auto tmp1 = std::vector<double>(rng1.size());
    std::partial_sum(rng1.begin(), rng1.end(), tmp1.begin());

    auto tmp2 = std::vector<double>(rng2.size());
    std::adjacent_difference(rng2.begin(), rng2.end(), tmp2.begin(),
                             [] (const auto left, const auto right)
    {
        return 3.0 * left - 2.0 * right;
    });

    return std::inner_product(tmp1.begin() + 1, tmp1.end(), tmp2.begin() + 1, 0.0);
}
```

# Optimizing, C++17

```
auto random_17_3(const std::vector<double>& rng1,
                 const std::vector<double>& rng2)
{
    double partial_sum = rng1[0];
    size_t i = 0;
    return std::accumulate(rng1.begin() + 1, rng1.end(), 0.0,
                           [&](const auto sum, const auto cur)
                           {
                               partial_sum += cur;
                               ++i;
                               return sum + prev * (3.0 * rng2[i] - 2.0 * rng2[i - 1]);
                           });
}
```



# Optimizing, C++23

```
double random_23(const std::vector<double>& rng)
{
    auto rng1 = random_23_1(rng);
    auto rng2 = random_23_2(rng1);
    return random_23_3(rng1, rng2);
}
```

```
auto random_23_1(const std::vector<double>& rng)
{
    return ranges::views::cartesian_product(
        rng,
        ranges::views::iota(1, 5),
        rng | ranges::views::transform([](const auto e) { return e * e + 2; })
        | ranges::views::transform([](const auto t)
        {
            const auto& [a, b, c] = t;
            return a * b + 2 * c;
        }));
}
```

# Optimizing, C++23

```
double random_23(const std::vector<double>& rng)
{
    auto rng1 = random_23_1(rng);
    auto rng2 = random_23_2(rng1);
    return random_23_3(rng1, rng2);
}
```

```
auto random_23_1(const std::vector<double>& rng)
{
    return ranges::views::cartesian_product(
        rng,
        ranges::views::iota(1, 5),
        rng | ranges::views::transform([](const auto e) { return e * e + 2; })
        | ranges::views::transform([](const auto t)
        {
            const auto& [a, b, c] = t;
            return a * b + 2 * c;
        }) | ranges::to<std::vector>;
}
```





# Google Benchmark



# Optimized Benchmarks

Run on (24 X 3450.41 MHz CPU s)

CPU Caches:

L1 Data 48 KiB (x12)

L1 Instruction 32 KiB (x12)

L2 Unified 2048 KiB (x12)

L3 Unified 30720 KiB (x1)

Benchmark	Time	CPU	Iterations
c_style_benchmark	25.3 ms	19.8 ms	34
cpp_17_benchmark	35.4 ms	23.8 ms	50
cpp_23_benchmark	317.0 ms	223.0 ms	4
c_style_benchmark_optimized	13.2 ms	9.58 ms	75
cpp_17_benchmark_optimized	18.2 ms	15.5 ms	90
cpp_23_benchmark_optimized	12.0 ms	6.20 ms	204



Cachegrind



# Cachegrind, optimized C – style C++

I refs: 76,367,324  
I1 misses: 2,025  
LLi misses: 1,956  
I1 miss rate: 0.00%  
LLi miss rate: 0.00%

D refs: 16,758,297 (12,556,469 rd + 4,201,828 wr)  
D1 misses: 2,515,039 ( 1,512,505 rd + 1,002,534 wr)  
LLd misses: 2,509,755 ( 1,507,993 rd + 1,001,762 wr)  
D1 miss rate: 15.0% ( 12.0% + 23.9% )  
LLd miss rate: 15.0% ( 12.0% + 23.8% )

LL refs: 2,517,064 ( 1,514,530 rd + 1,002,534 wr)  
LL misses: 2,511,711 ( 1,509,949 rd + 1,001,762 wr)  
LL miss rate: 2.7% ( 1.7% + 23.8% )

Branches: 5,334,767 ( 5,329,349 cond + 5,418 ind)  
Mispredicts: 18,704 ( 18,136 cond + 568 ind)  
Mispred rate: 0.4% ( 0.3% + 10.5% )



# Cachegrind, optimized C – style C++

I	refs:	92,364,114	76,367,324
I1	misses:	2,025	2,025
LLi	misses:	1,956	1,956
I1	miss rate:	0.00%	0.00%
LLi	miss rate:	0.00%	0.00%
D	refs:	26,758,586	16,758,297
D1	misses:	4,515,098	2,515,039
LLd	misses:	4,509,814	2,509,755
D1	miss rate:	16.9%	15.0%
LLd	miss rate:	16.9%	15.0%
LL	refs:	4,517,123	2,517,064
LL	misses:	4,511,770	2,511,711
LL	miss rate:	3.8%	2.7%
Branches:		11,334,889	5,334,767
Mispredicts:		18,731	18,704
Mispred rate:		0.2%	0.4%

# Cachegrind, optimized C++17

I refs: 164,384,131  
I1 misses: 2,033  
LLi misses: 1,965  
I1 miss rate: 0.00%  
LLi miss rate: 0.00%

D refs: 84,759,055 (14,557,331 rd + 70,201,724 wr)  
D1 misses: 3,515,169 ( 1,512,636 rd + 2,002,533 wr)  
LLd misses: 3,131,856 ( 1,508,160 rd + 1,623,696 wr)  
D1 miss rate: 4.1% ( 10.4% + 2.9% )  
LLd miss rate: 3.7% ( 10.4% + 2.3% )

LL refs: 3,517,202 ( 1,514,669 rd + 2,002,533 wr)  
LL misses: 3,133,821 ( 1,510,125 rd + 1,623,696 wr)  
LL miss rate: 1.3% ( 0.8% + 2.3% )

Branches: 72,338,768 (72,333,349 cond + 5,419 ind)  
Mispredicts: 18,732 ( 18,162 cond + 570 ind)  
Mispred rate: 0.0% ( 0.0% + 10.5% )

# Cachegrind, optimized C++17

I	refs:	252,376,633	164,384,131
I1	misses:	2,044	2,033
LLi	misses:	1,976	1,965
I1	miss rate:	0.00%	0.00%
LLi	miss rate:	0.00%	0.00%
D	refs:	144,802,262	84,759,055
D1	misses:	6,015,098	3,515,169
LLd	misses:	5,979,557	3,131,856
D1	miss rate:	4.2%	4.1%
LLd	miss rate:	4.1%	3.7%
LL	refs:	6,017,142	3,517,202
LL	misses:	5,981,533	3,133,821
LL	miss rate:	1.5%	1.3%
Branches:	115,330,825	72,338,768	
Mispredicts:	18,725	18,732	
Mispred rate:	0.0%	0.0%	



# Cachegrind, optimized C++23

I refs: 150,312,446  
I1 misses: 2,045  
LLi misses: 1,977  
I1 miss rate: 0.00%  
LLi miss rate: 0.00%

D refs: 24,756,927 (20,555,184 rd + 4,201,743 wr)  
D1 misses: 1,015,022 ( 512,478 rd + 502,544 wr)  
LLd misses: 645,605 ( 143,833 rd + 501,772 wr)  
D1 miss rate: 4.1% ( 2.5% + 12.0% )  
LLd miss rate: 2.6% ( 0.7% + 11.9% )

LL refs: 1,017,067 ( 514,523 rd + 502,544 wr)  
LL misses: 647,582 ( 145,810 rd + 501,772 wr)  
LL miss rate: 0.4% ( 0.1% + 11.9% )

Branches: 16,325,629 (16,320,221 cond + 5,408 ind)  
Mispredicts: 18,668 ( 18,102 cond + 566 ind)  
Mispred rate: 0.1% ( 0.1% + 10.5% )

# Cachegrind, optimized C++23

I	refs:	542,456,846	150,312,446
I1	misses:	2,036	2,045
LLi	misses:	1,936	1,977
I1	miss rate:	0.00%	0.00%
LLi	miss rate:	0.00%	0.00%
D	refs:	232,780,702	24,756,927
D1	misses:	14,845	1,015,022
LLd	misses:	9,425	645,605
D1	miss rate:	0.0%	4.1%
LLd	miss rate:	0.0%	2.6%
LL	refs:	16,881	1,017,067
LL	misses:	11,361	647,582
LL	miss rate:	0.0%	0.4%
Branches:		36,340,555	16,325,629
Mispredicts:		19,671	18,668
Mispred rate:		0.1%	0.1%

# Cachegrind compare optimized

	C-Style	C++17	C++23
I refs:	76,367,324	164,384,131	150,312,446
I1 misses:	2,025	2,033	2,045
LLi misses:	1,956	1,965	1,977
I1 miss rate:	0.00%	0.00%	0.00%
LLi miss rate:	0.00%	0.00%	0.00%
D refs:	16,758,297	84,759,055	24,756,927
D1 misses:	2,515,039	3,515,169	1,015,022
LLd misses:	2,509,755	3,131,856	645,605
D1 miss rate:	15.0%	4.1%	4.1%
LLd miss rate:	15.0%	3.7%	2.6%
LL refs:	2,517,064	3,517,202	1,017,067
LL misses:	2,511,711	3,133,821	647,582
LL miss rate:	2.7%	1.3%	0.4%
Branches:	5,334,767	72,338,768	16,325,629
Mispredicts:	18,704	18,732	18,668
Mispred rate:	0.4%	0.0%	0.1%





Next?





# Conclusion





# Questions?

Tina Ulbrich

ROSEN Technology and Research Center



@\_yulivee\_



tina\_ulbrich@gmx.de



#include <C++>

tinaul